

Problem Statement:To predict the Rainfall based on the Various Features of DataSet

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.linear_model import LinearRegression
4 from sklearn import preprocessing,svm
5 from sklearn.model_selection import train_test_split
6 import matplotlib.pyplot as plt
7 import seaborn as sns
```

In [2]:

```
1 df=pd.read_csv(r"C:\Users\HP\Downloads\rainfall in india 1901-2015.csv")
2 df
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	J&F
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	130
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	150
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	150
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	200
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	200
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	100
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	100
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	600
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	600
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	100

4116 rows × 19 columns

Data Preprocessing

Type *Markdown* and LaTeX: α^2

In [3]:

```
1 df.head()
```

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	4
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	5
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	3

In [4]:

```
1 df.tail()
```

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	

In [5]:

```
1 df.isnull().any()
```

Out[5]:

```
SUBDIVISION    False
YEAR           False
JAN            True
FEB            True
MAR            True
APR            True
MAY            True
JUN            True
JUL            True
AUG            True
SEP            True
OCT            True
NOV            True
DEC            True
ANNUAL         True
Jan-Feb        True
Mar-May        True
Jun-Sep        True
Oct-Dec        True
dtype: bool
```

In [6]:

```
1 df.fillna(method='ffill',inplace=True)
```

In [7]:

```
1 df.isnull().sum()
```

Out[7]:

```
SUBDIVISION    0
YEAR           0
JAN            0
FEB            0
MAR            0
APR            0
MAY            0
JUN            0
JUL            0
AUG            0
SEP            0
OCT            0
NOV            0
DEC            0
ANNUAL         0
Jan-Feb        0
Mar-May        0
Jun-Sep        0
Oct-Dec        0
dtype: int64
```

In [8]:

```
1 df.describe()
```

Out[8]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000
mean	1958.218659	18.957240	21.823251	27.415379	43.160641	85.788994	230.567979	347.177235	290.23979
std	33.140898	33.576192	35.922602	47.045473	67.816588	123.220150	234.896056	269.321089	188.78563
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000	0.000000	0.00000
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.475000	175.900000	155.85000
50%	1958.000000	6.000000	6.700000	7.900000	15.700000	36.700000	138.900000	284.800000	259.40000
75%	1987.000000	22.200000	26.800000	31.400000	50.125000	97.400000	306.150000	418.325000	377.80000
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000	2362.800000	1664.60000

In [9]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4116 non-null   float64
3   FEB             4116 non-null   float64
4   MAR             4116 non-null   float64
5   APR             4116 non-null   float64
6   MAY             4116 non-null   float64
7   JUN             4116 non-null   float64
8   JUL             4116 non-null   float64
9   AUG             4116 non-null   float64
10  SEP             4116 non-null   float64
11  OCT             4116 non-null   float64
12  NOV             4116 non-null   float64
13  DEC             4116 non-null   float64
14  ANNUAL          4116 non-null   float64
15  Jan-Feb         4116 non-null   float64
16  Mar-May         4116 non-null   float64
17  Jun-Sep         4116 non-null   float64
18  Oct-Dec         4116 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [10]:

```
1 df.columns
```

Out[10]:

```
Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
      'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May',
      'Jun-Sep', 'Oct-Dec'],
      dtype='object')
```

In [11]:

```
1 df.shape
```

Out[11]:

```
(4116, 19)
```

In [12]:

```
1 df['ANNUAL'].value_counts()
```

Out[12]:

```
ANNUAL
790.5      4
770.3      4
1836.2     4
1024.6     4
1926.5     3
..
443.9      1
689.0      1
605.2      1
509.7      1
1642.9     1
Name: count, Length: 3712, dtype: int64
```

In [13]:

```
1 df['Jan-Feb'].value_counts()
```

Out[13]:

Jan-Feb
0.0 238
0.1 80
0.2 52
0.3 38
0.4 32
...
23.3 1
95.2 1
76.9 1
66.5 1
69.3 1
Name: count, Length: 1220, dtype: int64

In [14]:

```
1 df['Mar-May'].value_counts()
```

Out[14]:

Mar-May
0.0 29
0.1 13
0.3 11
8.3 11
11.5 10
..
246.3 1
248.1 1
151.3 1
249.5 1
223.9 1
Name: count, Length: 2262, dtype: int64

In [15]:

```
1 df['Jun-Sep'].value_counts()
```

Out[15]:

Jun-Sep
434.3 4
334.8 4
573.8 4
613.3 4
1082.3 3
..
301.6 1
380.9 1
409.3 1
229.4 1
958.5 1
Name: count, Length: 3683, dtype: int64

In [16]:

```
1 df['Oct-Dec'].value_counts()
```


Out[16]:

```
Oct-Dec
0.0      16
0.1      15
0.5      13
0.6      12
0.7      11
..
191.5     1
124.5     1
139.1     1
41.5      1
555.4     1
Name: count, Length: 2389, dtype: int64
```

Data Analysis

In [17]:

```
1 df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]
2 sns.heatmap(df.corr(),annot=True)
3 plt.show()
```



In [18]:

```
1 df.columns
```

Out[18]:

```
Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```

In [19]:

```
1 x=df[["FEB"]]
2 y=df["JAN"]
```

Linear Regression

In [20]:

```
1 from sklearn.model_selection import train_test_split
2 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

In [21]:

```
1 from sklearn.linear_model import LinearRegression
2 reg=LinearRegression()
3 reg.fit(X_train,y_train)
4 print(reg.intercept_)
5 coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
6 coeff_
```

9.6506666612303553

Out[21]:

	coefficient
FEB	0.442278

In [22]:

```
1 score=reg.score(X_test,y_test)
2 print(score)
```

0.1793580786264921

In [23]:


```
1 predictions=reg.predict(X_test)
```

In [24]:

```
1 plt.scatter(y_test,predictions)
```


Out[24]:

<matplotlib.collections.PathCollection at 0x1c0f02591d0>



In [25]:

```
1 df500=df[:][:500]
2 sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)
3 plt.show()
```



In [26]:

```
1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
2 reg.fit(X_train,y_train)
3 reg.fit(X_test,y_test)
```

Out[26]:


```
LinearRegression
LinearRegression()
```


In [27]:

```

1 y_pred=reg.predict(X_test)
2 plt.scatter(X_test,y_test,color='black')
3 plt.plot(X_test,y_pred,color='red')
4 plt.show()

```



In [28]:

```

1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model=LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred=model.predict(X_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2 Score:",r2)

```

R2 Score: 0.13822029797956226

Ridge Regression

In [29]:

```

1 from sklearn.linear_model import Lasso,Ridge
2 from sklearn.preprocessing import StandardScaler

```

In [30]:

```

1 features= df.columns[0:5]
2 target= df.columns[-5]

```

In [31]:

```

1 x=np.array(df['JAN']).reshape(-1,1)
2 y=np.array(df['FEB']).reshape(-1,2)

```

In [32]:

```

1 x= df[features].values
2 y= df[target].values
3 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)

```

In [33]:

```
1 ridgeReg=Ridge(alpha=10)
2 ridgeReg.fit(x_train,y_train)
3 train_score_ridge=ridgeReg.score(x_train,y_train)
4 test_score_ridge=ridgeReg.score(x_test,y_test)
```

In [34]:

```
1 print("\n Ridge Model:\n")
2 print("the train score for ridge model is{}".format(train_score_ridge))
3 print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

```
the train score for ridge model is0.9999999999874192
the test score for ridge model is0.99999999998833
```

In [35]:


```
1 lr=LinearRegression()
```

In [36]:

```

figure(figsize= (10,10))
plt.plot(Features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge; $\alpha = 10$')
plt.plot(Features,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green",label='Linear Regression')
plt.xticks(rotation = 90)
plt.grid()
plt.show()

```



Lasso Regression

In [37]:

```

1 print("\n Lasso Model:\n")
2 lasso=Lasso(alpha=10)
3 lasso.fit(x_train,y_train)
4 train_score_ls=lasso.score(x_train,y_train)
5 test_score_ls=lasso.score(x_test,y_test)
6 print("The train score for ls model is {}".format(train_score_ls))
7 print("The test score for ls model is {}".format(test_score_ls))

```

Lasso Model:


The train score for ls model is 0.9999207747038827
 The test score for ls model is 0.9999206791315255

In [38]:

```
1 pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

Out[38]:

<Axes: >



In [39]:

```
1 from sklearn.linear_model import LassoCV
2 lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_train)
3 print(lasso_cv.score(x_train,y_train))
4 print(lasso_cv.score(x_test,y_test))
```

0.9999999999999921


0.9999999999999921

In [40]:

```

figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge; $\alpha = 10$')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge; $\alpha = 100$')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation = 90)
plt.grid()
plt.show()

```



Elastic Net:

In [41]:

```

1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(x,y)
4 print(regr.coef_)
5 print(regr.intercept_)
6 print(regr.score(x,y))

```

```

[9.99098574e-01 0.00000000e+00 3.02728910e-05 0.00000000e+00
 0.00000000e+00]
0.016258606966612632
0.9999992160905338

```

In [42]:

```
1 y_pred_elastic = regr.predict(x_train)
```

In [43]:

```
1 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
2 print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 0.0008816302333951295

Conclusion:

In []:

given insurance data set have performed linear, logistic, random forest and decision tree models of regression and concluded that the most accuracy is occurred in logistic regression, i.e 99percent
e that the Logic Regression model is best fit for given data

In []:

```
1
```

In []:

```
1
```